

# Final code with MLP

April 19, 2018

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
        from sklearn.linear_model import LogisticRegression
        from sklearn.feature_selection import SelectKBest
        from sklearn.feature_selection import chi2
        from sklearn.metrics import roc_auc_score
        from sklearn.tree import DecisionTreeClassifier
        #deep
        from sklearn.neural_network import MLPClassifier
        from sklearn.neural_network import MLPRegressor

import datetime
import pandas as pd
import numpy as np
import time

def log(text, t_start=None):
    if t_start is None:
        print(text)
    else:
        elapsed_time = round(time.time() - t_start, 2)
        print(text + "\t(" + str(elapsed_time) + "s)")

t = time.time()
customers = pd.read_csv("data/customers.csv")
products = pd.read_csv("data/products.csv")
x_train = pd.read_csv("data/X_train.csv")
y_train = pd.read_csv("data/y_train.csv")
```

```

x_test = pd.read_csv("data/X_test.csv")
log("files loaded", t)

# SizeAdviceDescription
SizeAdviceDescriptionCleaner = {}
SizeAdviceDescriptionCleaner['nan'] = 0
SizeAdviceDescriptionCleaner['Ce mod\xc3\x83\xc2\xa8le chausse normalement'] = 0
SizeAdviceDescriptionCleaner['Mod\xc3\x83\xc2\xa8le confortable, convient aux pieds larges'] = -.5
SizeAdviceDescriptionCleaner['Mod\xc3\x83\xc2\xa8le \xc3\x83\xc2\xa9troit, convient aux pieds fins'] = .5
SizeAdviceDescriptionCleaner['Prenez votre pointure habituelle'] = 0
SizeAdviceDescriptionCleaner['Chaussant particul\xc3\x83\xc2\xa8rement g\xc3\x83\xc2\xa9n\xc3\x83\xc2\xa9reux. Nous vous conseil']
SizeAdviceDescriptionCleaner['Chaussant petit. Si vous \xc3\x83\xc2\xaates habituellement entre deux pointures, nous vous conseil']
SizeAdviceDescriptionCleaner['Prenez une taille au-dessus de sa pointure !'] = 1
SizeAdviceDescriptionCleaner['Prenez une taille au-dessus de votre pointure habituelle'] = 1
SizeAdviceDescriptionCleaner['Prenez une taille en dessous de sa pointure !'] = -1
SizeAdviceDescriptionCleaner['Prenez une taille en dessous de votre pointure habituelle'] = -1

# BirthDate
def age(birthdate):
    if type(birthdate) == type(" "):
        return 2016 - int(birthdate[:4])
    return None

# OrderCreationDate and SeasonLabel
def order_season(orderdate):
    month = int(orderdate[5:7])
    if month >= 4 and month <= 9:
        return "Printemps/Et\xc3\x83\xc2\xa9"
    return "Automne/Hiver"

def build_df(x):
    """Builds a pandas DataFrame with clean columns from a read CSV"""

    t = time.time()
    m = None

    # join
    m = pd.merge(x, products, how='left', on='VariantId', suffixes=('_pr', ''))

```

```

m = pd.merge(m, customers, how='left', on='CustomerId', suffixes=('_cs', ''))

# converting UnitPMPEUR
m.UnitPMPEUR = m["UnitPMPEUR"].map(lambda row: float(row.replace(',', ' ')))

# building news columns
m["MatchGender"] = m["Gender"] == m["GenderLabel"]
m["MatchSeason"] = m["SeasonLabel_pr"] == m["SeasonLabel"]
m["OrderSeason"] = m["OrderCreationDate"].map(order_season)
m["MatchOrderSeason"] = m["OrderSeason"] == m["SeasonLabel"]

# cleaning
m["SizeAdviceDescription"] = m["SizeAdviceDescription"].map(SizeAdviceDescriptionCleaner)
m["BirthDate"] = m["BirthDate"].map(age)

# removing useless columns
blacklist = ['VariantId', 'CustomerId', 'OrderNumber', 'LineItem',
             'ProductColorId', 'BrandId', 'SupplierColor', 'OrderShipDate',
             'ProductId', 'BillingPostalCode', 'FirstOrderDate',
             'OrderStatusLabel', 'MinSize', 'MaxSize', 'OrderSeason',
             'OrderCreationDate', 'SubtypeLabel', 'ProductType']
whitelist = None
if blacklist is not None:
    m = m.drop(blacklist, axis=1)
if whitelist is not None:
    for col in m.columns:
        if col not in whitelist:
            m = m.drop([col], axis=1)

print "dataframe shape:", m.shape
log("dataframe built", t)
return m

```

```

df_test = build_df(x_test)
df_train = build_df(x_train)

```

```

def mask(m):
    columns2bin = [col for col in m.columns if m[col].dtype == 'object']
    other_cols = m.drop(columns2bin, axis=1)
    new_cols = pd.get_dummies(m.loc[:, columns2bin])
    res = pd.concat([other_cols, new_cols], axis=1)
    res = res.fillna(0)
    print "new shape:", res.shape
    return res

def compute(name, clf, x1, x2, slc=100000):
    print "\n-----", name, "-----"
    clf.fit(x1.iloc[:slc], y_train.ReturnQuantityBin[:slc])

    predict_train = clf.predict_proba(x1.iloc[:slc])
    score_train = roc_auc_score(y_train.ReturnQuantityBin[:slc], predict_train[:, 1])
    print "train score:", score_train

    predict_test = clf.predict_proba(x1.iloc[slc:2 * slc])
    score_test = roc_auc_score(y_train.ReturnQuantityBin[slc:2 * slc], predict_test[:, 1])
    print "test score:", score_test
    return score_train, score_test

def compute_all(x1, x2, slc=100000):
    """Tries different classifiers and returns the best one (best test score)"""
    t = time.time()
    best_index, best_score = None, None

    print "train shape:\t", x1.shape, "\t", y_train.shape
    print "test shape:\t", x2.shape, "\t", y_test.shape

    classifiers = [
        ("random forest", RandomForestClassifier()),
        ("decision tree", DecisionTreeClassifier()),
        ("logistic regression", LogisticRegression()),
        ("DEEP", MLPClassifier(solver='lbfgs', alpha=1e-5,
                               hidden_layer_sizes=(100, 10), random_state=1))
    ]

    for i, (name, clf) in enumerate(classifiers):
        score_train, score_test = compute(name, clf, x1, x2, slc)
        if best_score is None or score_test > best_score:

```

```

        best_index, best_score = i, score_test

    log("\nbest classifier: " + classifiers[best_index][0], t)
    return classifiers[best_index][1]

def output(clf, x1, x2):
    t = time.time()
    y_tosubmit = clf.predict_proba(x2.loc[:, x1.columns].fillna(0))

    timestamp = '{0:%Y_%m_%d_%H_%M_%S}'.format(datetime.datetime.now())
    filename = "ypred_{0}_sgd.txt".format(timestamp)
    np.savetxt(filename, y_tosubmit[:,1], fmt='%f')

    f = open("predictions.txt", 'a')
    f.write(timestamp + '\n' + repr(clf).replace('\n', ' ') + '\n\n')
    f.close()

    print "shape:", y_tosubmit.shape
    log("generated output at " + filename, t)

t = time.time()
x1 = mask(df_train)
x2 = mask(df_test)
log("applied mask", t)

def shuffle(x, y, steps=10, slc=100000, plot=True):
    scores_train, scores_test = [], []
    best_clf, best_score = None, None

    z = x.copy(deep=True)
    z["ReturnQuantityBin"] = y.ReturnQuantityBin

    for k in range(steps):
        u = z.sample(frac=1)
        v = u.loc[:, ["ReturnQuantityBin"]]
        u = u.drop(["ReturnQuantityBin"], axis=1)

```

```

#clf = LogisticRegression()
clf = MLPRegressor(solver='sgd', alpha=1e-5,
                  hidden_layer_sizes=(100, 100,100,100), random_state=1)
clf.fit(u.iloc[:slc], v.ReturnQuantityBin[:slc])

predict_train = clf.predict_proba(u.iloc[:slc])
score_train = roc_auc_score(v.ReturnQuantityBin[:slc], predict_train[:, 1])

predict_test = clf.predict_proba(u.iloc[slc:2 * slc])
score_test = roc_auc_score(v.ReturnQuantityBin[slc:2 * slc], predict_test[:, 1])

if best_clf is None or score_test > best_score:
    best_clf, best_score = clf, score_test

if plot:
    print "test", k, "\tttrain:", score_train, "\tttest:", score_test

scores_train.append(score_train)
scores_test.append(score_test)

if plot:
    plt.figure(figsize=(16, 10))
    plt.xlabel("train score")
    plt.ylabel("test score")
    plt.plot(scores_train, scores_test, '+')
    plt.show()

return scores_train, scores_test, best_clf, best_score

sc_train, sc_test, clf, score = shuffle(x1, y_train, slc=100000, steps=10, plot=False)

output(clf, x1, x2)

```

# Code with all trials

April 19, 2018

## 1 AlphaShoe

SD210 Challenge

- [scoreboard and submissions](#)
- [starting kit](#)
- [training data](#)
- [dictionnary](#)

### 1.0.1 Importations

In [1]: *# coding: utf-8*

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from sklearn.metrics import roc_auc_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder
from matplotlib import pyplot as plt
import datetime
import pandas as pd
import numpy as np
import time
```

```
%matplotlib inline

def log(text, t_start=None):
    if t_start is None:
        print(text)
    else:
        elapsed_time = round(time.time() - t_start, 2)
        print(text + "\t(" + str(elapsed_time) + "s)")
```

## 1.0.2 Loading files

```
In [2]: t = time.time()
        customers = pd.read_csv("data/customers.csv")
        products = pd.read_csv("data/products.csv")
        x_train = pd.read_csv("data/X_train.csv")
        y_train = pd.read_csv("data/y_train.csv")
        x_test = pd.read_csv("data/X_test.csv")
        y_test = pd.read_csv("data/y_test.csv")
        log("files loaded", t)
```

files loaded (43.61s)

## 1.0.3 Data cleaning

```
In [3]: # SizeAdviceDescription
        SizeAdviceDescriptionCleaner = {}
        SizeAdviceDescriptionCleaner['nan'] = 0
        SizeAdviceDescriptionCleaner['Ce mod\&#x83\&#xc2\&#xa8le chausse normalement'] = 0
        SizeAdviceDescriptionCleaner['Mod\&#x83\&#xc2\&#xa8le confortable, convient aux pieds larges'] = -.5
        SizeAdviceDescriptionCleaner['Mod\&#x83\&#xc2\&#xa8le \&#x83\&#xc2\&#xa9troit, convient aux pieds fins'] = .5
        SizeAdviceDescriptionCleaner['Prenez votre pointure habituelle'] = 0
        SizeAdviceDescriptionCleaner['Chaussant particul\&#x83\&#xc2\&#xa8rement g\&#x83\&#xc2\&#xa9n\&#x83\&#xc2\&#xa9reux. Nous vous conseil
        SizeAdviceDescriptionCleaner['Chaussant petit. Si vous \&#x83\&#xc2\&#xaates habituellement entre deux pointures, nous vous conseil
        SizeAdviceDescriptionCleaner['Prenez une taille au-dessus de sa pointure !'] = 1
        SizeAdviceDescriptionCleaner['Prenez une taille au-dessus de votre pointure habituelle'] = 1
        SizeAdviceDescriptionCleaner['Prenez une taille en dessous de sa pointure !'] = -1
        SizeAdviceDescriptionCleaner['Prenez une taille en dessous de votre pointure habituelle'] = -1
```



```

# datetime format:
# YYYY-MM-DD HH:mm:ss

# BirthDate
def age(birthdate):
    if type(birthdate) == type(" "):
        return 2016 - int(birthdate[:4])
    return None

def week_of_the_year(date):
    if type(date) == type(" "):
        year = int(date[:4])
        month = int(date[5:7])
        day = int(date[8:10])
        return datetime.date(year, month, day).isocalendar()[1]
    return None

def hour(date):
    if type(date) == type(" "):
        return int(date[11:13])
    return None

# OrderCreationDate and SeasonLabel
def order_season(orderdate):
    month = int(orderdate[5:7])
    if month >= 4 and month <= 9:
        return "Printemps/Et\x83\x83\x82\xa9"
    return "Automne/Hiver"

def build_df(x):
    """Builds a pandas DataFrame with clean columns from a read CSV"""

    t = time.time()
    m = None

    # join
    m = pd.merge(x, products, how='left', on='VariantId', suffixes=('_pr', ''))
    m = pd.merge(m, customers, how='left', on='CustomerId', suffixes=('_cs', ''))

```

```

# converting UnitPMPEUR
m.UnitPMPEUR = m["UnitPMPEUR"].map(lambda row: float(row.replace(',', '.)))

# building news columns
m["MatchGender"] = m["Gender"] == m["GenderLabel"]
m["MatchSeason"] = m["SeasonLabel_pr"] == m["SeasonLabel"]
m["OrderSeason"] = m["OrderCreationDate"].map(order_season)
m["MatchOrderSeason"] = m["OrderSeason"] == m["SeasonLabel"]
m["OrderHour"] = m["OrderCreationDate"].map(hour)
m["FirstOrderDate"] = m["FirstOrderDate"].map(week_of_the_year)

# cleaning
m["SizeAdviceDescription"] = m["SizeAdviceDescription"].map(SizeAdviceDescriptionCleaner)
m["BirthDate"] = m["BirthDate"].map(age)

# removing useless columns
blacklist = ['VariantId', 'CustomerId', 'OrderNumber', 'LineItem',
             'ProductColorId', 'OrderShipDate',
             'ProductId', 'BillingPostalCode', 'SupplierColor',
             'OrderStatusLabel',
             'OrderCreationDate', 'SubtypeLabel', 'ProductType',
             ]
whitelist = None
if blacklist is not None:
    m = m.drop(blacklist, axis=1)
if whitelist is not None:
    for col in m.columns:
        if col not in whitelist:
            m = m.drop([col], axis=1)

print "dataframe shape:", m.shape
log("dataframe built", t)
return m

```

#### 1.0.4 Statistics

```

In [21]: def returns_frequency(x, y, col, step):
         """Returns the returns frequencies for each value of a column"""

```

```

counter = 0

# counting occurrences of each column value
occurrences = {}
for i, o in x.loc[:, :step].iterrows():
    counter += 1
    if str(o[col]) not in occurrences.keys():
        occurrences[str(o[col])] = [0., 0.]
    if y.loc[i, ["ReturnQuantityBin"]][0] == 0.0:
        occurrences[str(o[col])][0] += 1.
    else:
        occurrences[str(o[col])][1] += 1.

# computing the returns frequency, stored in `recap`
recap, values = [], []
for val, (zeros, ones) in occurrences.items():
    values.append(((ones / (zeros + ones))))
    recap.append((val, values[-1]))
recap.sort(key=lambda row: row[0])

# computing variance and relative variance
var = np.var(values)
rel = var / len(values)

return recap, var, rel, counter

def column_stats(x, y, col, step, verbose):
    """Computes the statistics for one column"""

    print "\n----- " + col + " -----"
    recap, var, rel, counter = returns_frequency(x, y, col, step)

    if verbose:
        for (val, freq) in recap:
            print val, "\t", freq, "returns"

    print "variance:", round(var, 5), "\tvalues count:", len(recap)
    return recap, var, rel, counter

```

```

def compute_statistics(x, y, blacklist=[], whitelist=[], step=100, verbose=False):
    ignored = []
    labels_g, labels_n = [], []
    scattering_g, scattering_n = [], []      # variances
    r_scattering_g, r_scattering_n = [], []  # variance divided by the number of different values

    counter_cols = 0
    counter_rows = 0

    for col in x.columns:
        if col not in blacklist and x[col].dtype in ["object", "bool"]:
            counter_cols += 1
            labels_g.append(col)
            recap, var, rel, counter_rows = column_stats(x, y, col, step, verbose)
            scattering_g.append(var)
            r_scattering_g.append(rel)
        elif col not in blacklist and x[col].dtype in ["float64", "int64"]:
            counter_cols += 1
            labels_n.append(col)
            recap, var, rel, counter_rows = column_stats(x, y, col, step, verbose)
            scattering_n.append(var)
            r_scattering_n.append(rel)
            if col in whitelist:
                plot_dots([float(v[0]) for v in recap], [v[1] for v in recap], col)
        else:
            ignored.append(col)

    print "\n\nanalyzed", counter_cols, "columns and", counter_rows, "rows"
    print "\n\nignored columns:", ignored
    plot_barchart(labels_g, scattering_g, r_scattering_g, "general columns")
    plot_barchart(labels_n, scattering_n, r_scattering_n, "numerical columns")

def plot_barchart(labels, values_l, values_r, title):
    fig, ax = plt.subplots(figsize=(16, 10))
    ind = np.arange(len(values_l))
    width = .35

```

```

ax.bar(ind - width/2, values_l, width)
ax.bar(ind + width/2, values_r, width)
ax.set_xticks(ind)
ax.set_xticklabels(labels)
plt.xticks(rotation=70)
plt.title(title)
plt.show()

```

```

def plot_dots(xs, ys, xlabel):
    plt.figure(figsize=(8, 5))
    plt.xlabel(xlabel)
    plt.ylabel("returns frequency")
    plt.plot(xs, ys, 'o')
    plt.show()

```

```
In [22]: df_stats = build_df(x_train)
```

```

dataframe shape: (1067290, 40)
dataframe built      (76.23s)

```

```
In [6]: compute_statistics(df_stats, y_train, blacklist=[], whitelist=df_stats.columns, step=100)
```

```

----- OrderTypelabel -----
variance: 0.00017          values count: 2

----- SeasonLabel_pr -----
variance: 7e-05           values count: 2

----- PaymentModelLabel -----
variance: 0.00472          values count: 10

----- CustomerTypeLabel -----
variance: 0.00022          values count: 2


----- IsoCode -----
variance: 0.04536          values count: 20

```


----- DeviceTypeLabel -----  
variance: 0.00124            values count: 4

----- PricingTypeLabel -----  
variance: 0.00072            values count: 5


----- TotalLineItems -----  
variance: 0.13073            values count: 26



----- Quantity -----  
variance: 0.01033      values count: 4




----- UnitPMPEUR -----  
variance: 0.13264      values count: 5066




----- OrderNumCustomer -----  
variance: 0.12625      values count: 67





```
----- IsOnSale -----  
variance: 0.00176      values count: 3
```



----- GenderLabel -----

variance: 0.00338          values count: 6


----- MarketTargetLabel -----

variance: 0.00735          values count: 14


----- SeasonLabel -----

variance: 0.00025          values count: 3

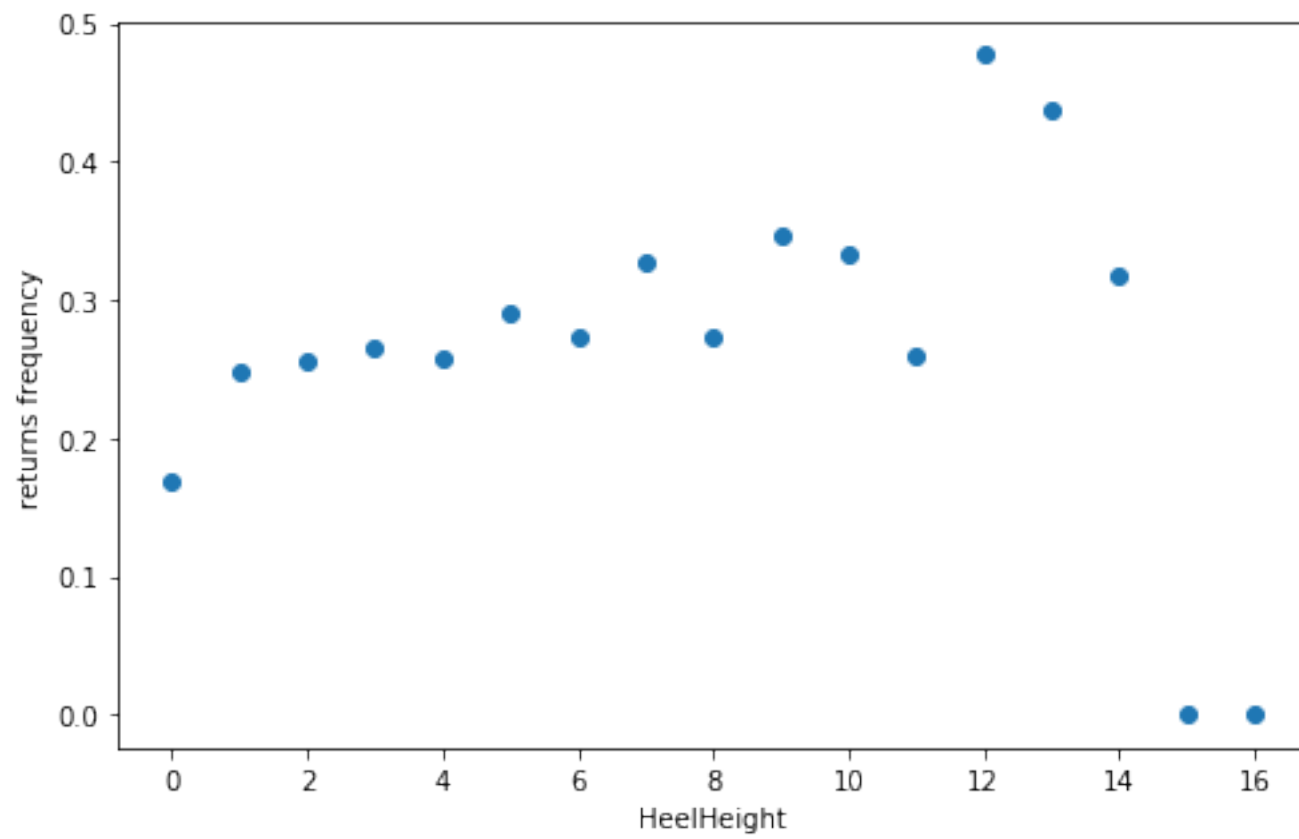
```
----- SeasonalityLabel -----  
variance: 0.00057          values count: 4  
  
----- UniverseLabel -----  
variance: 0.00268          values count: 9  
  
----- TypeBrand -----  
variance: 0.00581          values count: 4  
  
----- CalfTurn -----  
variance: 0.08768          values count: 20
```




----- UpperHeight -----  
variance: 0.05583      values count: 49




----- HeelHeight -----  
variance: 0.01392      values count: 18



----- PurchasePriceHT -----  
variance: 0.11134      values count: 1446



```
----- IsNewCollection -----
variance: 0.00217      values count: 3
```



```
----- UpperMaterialLabel -----  
variance: 0.00883      values count: 9
```


```
----- LiningMaterialLabel -----  
variance: 0.07434      values count: 11
```

```
----- OutSoleMaterialLabel -----  
variance: 0.01128      values count: 8
```



----- RemovableSole -----  
variance: 0.00039          values count: 3


----- SizeAdviceDescription -----  
variance: 0.0099          values count: 7



----- CountryISOCode -----

variance: 0.05302      values count: 19

----- BirthDate -----  
variance: 0.03355      values count: 87



----- Gender -----  
variance: 0.00061      values count: 3


```
----- MatchGender -----  
variance: 0.00048          values count: 2
```


```
----- MatchSeason -----  
variance: 0.00032          values count: 2
```

```
----- MatchOrderSeason -----  
variance: 0.00017          values count: 2
```

analyzed 34 columns and 10673 rows

ignored columns: []





```
In [26]: column_stats(df_train, y_train, "Gender", 100, True)
```

```
----- Gender -----  
Femme          0.212231857599 returns  
Homme           0.160953800298 returns  
nan            0.214285714286 returns  
variance: 0.00061          values count: 3
```

```
Out[26]: ([('Femme', 0.21223185759926974),  
          ('Homme', 0.16095380029806258),  
          ('nan', 0.21428571428571427)],  
          0.0006086611708289573,  
          0.0002028870569429858,  
          10673)
```

### 1.0.5 Classification

```
In [28]: def mask_bin(m):  
    columns2bin = [col for col in m.columns if m[col].dtype == 'object']  
    other_cols = m.drop(columns2bin, axis=1)  
    new_cols = pd.get_dummies(m.loc[:, columns2bin])  
    res = pd.concat([other_cols, new_cols], axis=1)  
    res = res.fillna(0)  
    print "new shape:", res.shape  
    return res  
  
def mask_int(m):  
    columns2int = [col for col in m.columns if m[col].dtype == 'object']  
    res = m.apply(LabelEncoder().fit_transform)  
    res = res.fillna(0)  
    print "new shape:", res.shape  
    return res  
  
def compute(name, clf, x1, x2, slc=100000):  
    print "\n-----", name, "-----"  
    clf.fit(x1.iloc[:slc], y_train.ReturnQuantityBin[:slc])  
  
    predict_train = clf.predict_proba(x1.iloc[:slc])
```

```

score_train = roc_auc_score(y_train.ReturnQuantityBin[:slc], predict_train[:, 1])
print "train score:", score_train

predict_test = clf.predict_proba(x1.iloc[slc:2 * slc])
score_test = roc_auc_score(y_train.ReturnQuantityBin[slc:2 * slc], predict_test[:, 1])
print "test score:", score_test
return score_train, score_test

def compute_all(x1, x2, slc=100000):
    """Tries different classifiers and returns the best one (best test score)"""
    t = time.time()
    best_index, best_score = None, None

    print "train shape:\t", x1.shape, "\t", y_train.shape
    print "test shape:\t", x2.shape, "\t", y_test.shape

    classifiers = [("random forest", RandomForestClassifier()),
                   ("decision tree", DecisionTreeClassifier()),
                   ("logistic regression", LogisticRegression())]

    for i, (name, clf) in enumerate(classifiers):
        score_train, score_test = compute(name, clf, x1, x2, slc)
        if best_score is None or score_test > best_score:
            best_index, best_score = i, score_test

    log("\nbest classifier: " + classifiers[best_index][0], t)
    return classifiers[best_index][1]

def output(clf, x1, x2):
    t = time.time()
    y_tosubmit = clf.predict_proba(x2.loc[:, x1.columns].fillna(0))

    timestamp = '{0:%Y_%m_%d_%H_%M_%S}'.format(datetime.datetime.now())
    filename = "ypred_{0}.txt".format(timestamp)
    np.savetxt(filename, y_tosubmit[:, 1], fmt='%f')

    f = open("predictions.txt", 'a')
    f.write(timestamp + '\n' + repr(clf).replace('\n', ' ') + '\n\n')
    f.close()

```

```
print "shape:", y_tosubmit.shape
log("generated output at " + filename, t)
```

### 1.0.6 Computation test loop

```
In [6]: df_train = build_df(x_train)
```

```
dataframe shape: (1067290, 40)
dataframe built      (6.42s)
```

```
In [27]: df_test = build_df(x_test)
```

```
dataframe shape: (800468, 40)
dataframe built      (11.46s)
```

```
In [14]: t = time.time()
         x1 = mask_bin(df_train)
         x2 = mask_bin(df_test)
         log("applied mask", t)
```

```
new shape: (1067290, 40)
new shape: (800468, 40)
applied mask      (24.57s)
```

```
In [9]: clf = compute_all(x1, x2)
```

```
train shape:      (1067290, 158)      (1067290, 4)
test shape:      (800468, 169)      (800468, 4)
```

```
----- random forest -----
train score: 0.9949450061129895
test score: 0.5720997532776027
```

```
----- decision tree -----
train score: 0.9993978968884667
test score: 0.5333491459838888
```



```
----- logistic regression -----  
train score: 0.63903864972542  
test score: 0.6329470043824317
```

```
best classifier: logistic regression      (14.0s)
```

```
In [13]: output(clf, x1, x2)
```

```
shape: (800468, 2)  
generated output at ypred_2018_04_05_15_54_25.txt      (5.94s)
```

### 1.0.7 Dataset permutation

```
In [ ]: def shuffle(x, y, clf_fun, steps=10, slc=50000, plot=True):  
    scores_train, scores_test = [], []  
    best_clf, best_score = None, None  
  
    z = x.copy(deep=True)  
    z["ReturnQuantityBin"] = y.ReturnQuantityBin  
  
    for k in range(steps):  
        u = z.sample(frac=1)  
        v = u.loc[:, ["ReturnQuantityBin"]]  
        u = u.drop(["ReturnQuantityBin"], axis=1)  
  
        clf = clf_fun()  
        clf.fit(u.iloc[:slc], v.ReturnQuantityBin[:slc])  
  
        predict_train = clf.predict_proba(u.iloc[:slc])  
        score_train = roc_auc_score(v.ReturnQuantityBin[:slc], predict_train[:, 1])  
  
        predict_test = clf.predict_proba(u.iloc[slc:2 * slc])  
        score_test = roc_auc_score(v.ReturnQuantityBin[slc:2 * slc], predict_test[:, 1])  
  
        if best_clf is None or score_test > best_score:  
            best_clf, best_score = clf, score_test  
  
    if plot:
```

```

        print "test", k, "\ttrain:", score_train, "\ttest:", score_test

    scores_train.append(score_train)
    scores_test.append(score_test)

if plot:
    plt.figure(figsize=(16, 10))
    plt.xlabel("train score")
    plt.ylabel("test score")
    plt.plot(scores_train, scores_test, '+')
    plt.show()

return scores_train, scores_test, best_clf, best_score

def try_slice(x, y, clf_fun, plot=True, steps=10):

    slices, scores, classifiers = [], [], {}
    best_slice, best_score = None, None

    for slc in range(10000, 200001, 10000):
        sc_train, sc_test, clf, score = shuffle(x, y, clf_fun, slc=slc, steps=steps, plot=False)
        slices.append(slc)
        scores.append(score)
        if best_slice is None or best_scores < score:
            best_slice, best_score = slc, score
        classifiers[slc] = clf
        print "slice", slc, "\t", score

    if plot:
        plt.figure(figsize=(16, 10))
        plt.xlabel("slice")
        plt.ylabel("score")
        plt.plot(slices, scores, '-o')
        plt.show()

    return classifiers[best_slice]

In [ ]: sc_train, sc_test, best_clf, best_score = shuffle(x1, y_train, slc=130000, steps=10, plot=False)
        print best_score

```

```
In [ ]: output(best_clf, x1, x2)
```

### 1.0.8 Principal component analysis (PCA)

Best scores reached with `n_components` at 96: - train score: 0.6388630967451936 - test score: 0.6331956289779485

Above, scores are decreasing.

**Note:** the result from `.transform()` is a Numpy array. Therefore the slicing is different.

```
In [29]: def try_pca(data, n_components):
    print "\n----- PCA", n_components, "-----"

    pca = PCA(n_components=n_components)
    pca.fit(data)
    x = pca.transform(data)
    clf = LogisticRegression()
    slc = 100000
    clf.fit(x[:slc, :], y_train.ReturnQuantityBin[:slc])

    predict_train = clf.predict_proba(x[:slc, :])
    score_train = roc_auc_score(y_train.ReturnQuantityBin[:slc], predict_train[:, 1])
    print "train score:", score_train

    predict_test = clf.predict_proba(x[slc:2 * slc, :])
    score_test = roc_auc_score(y_train.ReturnQuantityBin[slc:2 * slc], predict_test[:, 1])
    print "test score:", score_test

    return score_train, score_test, pca

for n in range(1, 100, 5):
    try_pca(x1, n)
```

```
----- PCA 1 -----
train score: 0.5702366075653182
test score: 0.5647601895068738
```

```
----- PCA 6 -----
train score: 0.5999356153282647
test score: 0.5983282609508912
```

```
----- PCA 11 -----  
train score: 0.6218118708758879  
test score: 0.6138734988058648  
  
----- PCA 16 -----  
train score: 0.6277049263505473  
test score: 0.6225947759354865  
  
----- PCA 21 -----  
train score: 0.6284750476572429  
test score: 0.6244283257910828  
  
----- PCA 26 -----  
train score: 0.6293423130492306  
test score: 0.6249607267334103  
  
----- PCA 31 -----  
train score: 0.6300006665601142  
test score: 0.6252857562462173  
  
----- PCA 36 -----  
train score: 0.6341085795216018  
test score: 0.6283580092583988  
  
----- PCA 41 -----  
train score: 0.6370415100553493  
test score: 0.6315438780860905  
  
----- PCA 46 -----  
train score: 0.6369463535588813  
test score: 0.6318376492846729  
  
----- PCA 51 -----  
train score: 0.6373923821818226  
test score: 0.632345993128418  
  
----- PCA 56 -----  
train score: 0.6376517894809209  
test score: 0.6320991622328517
```

```

----- PCA 61 -----
train score: 0.6377169758518486
test score: 0.6325997920070896

----- PCA 66 -----
train score: 0.638039089427804
test score: 0.63284863566079

----- PCA 71 -----
train score: 0.6384357857447727
test score: 0.6332645079580541

----- PCA 76 -----
train score: 0.6387240315583405
test score: 0.6331850942121828

----- PCA 81 -----
train score: 0.6387008959275011
test score: 0.6330236438208361

----- PCA 86 -----
train score: 0.6388413465503924
test score: 0.6331544222007438

----- PCA 91 -----
train score: 0.6388561386698401
test score: 0.6331928959049634

----- PCA 96 -----
train score: 0.6388630967451936
test score: 0.6331956289779485

```

### 1.0.9 Cross validation and random forests

This is an *attempt*, scores did not met expectation.

```

In [61]: def cross_val(name, clf, slc=100000):
          print "\n-----", name, "-----"

```

```

    predict_train = cross_val_predict(clf, x1.iloc[:slc], y_train.ReturnQuantityBin[:slc], cv=10, method='predict_proba')
    score_train = roc_auc_score(y_train.ReturnQuantityBin[:slc], predict_train[:, 1])
    print "train score:", score_train
    return score_train

cross_val("LogisticRegression", LogisticRegression())

```

```

----- LogisticRegression -----
train score: 0.6332445684201241

```

```
Out[61]: 0.6332445684201241
```

```
In [20]: depths, scores = [], []
```

```

for max_depth in range(10, 201, 10):
    clf = RandomForestClassifier(max_depth=max_depth)
    score = cross_val_score(clf, x1.loc[:50000], y_train.loc[:50000, "ReturnQuantityBin"], cv=5).mean()
    depths.append(max_depth)
    scores.append(score)
    print 'max_depth:', max_depth, '\tscore:', score

```


```

max_depth: 10      score: 0.7904241907464419
max_depth: 20      score: 0.7837043827060438
max_depth: 30      score: 0.7641646465896466
max_depth: 40      score: 0.7519450905184509
max_depth: 50      score: 0.7520049165158491
max_depth: 60      score: 0.7518849505168496
max_depth: 70      score: 0.7509051305164514
max_depth: 80      score: 0.7523450245226503
max_depth: 90      score: 0.7530050525222506
max_depth: 100     score: 0.7518051165306512
max_depth: 110     score: 0.7562650105484501
max_depth: 120     score: 0.7516650545148506
max_depth: 130     score: 0.7531449625248496
max_depth: 140     score: 0.7515450645214508
max_depth: 150     score: 0.7527851165222512
max_depth: 160     score: 0.7530050285102503
max_depth: 170     score: 0.7541649545348494

```

```
max_depth: 180          score: 0.7515251585076517
max_depth: 190          score: 0.7510651385024513
max_depth: 200          score: 0.7521451345096513
```

```
In [21]: plt.figure(figsize=(16, 10))
         plt.xlabel("depth")
         plt.ylabel("test score")
         plt.plot(depths, scores, '-+')
         plt.show()
```



### 1.0.10 Feature importance

First, we try to find an *optimal maximum* depth for a random forest classifier using cross validation.

In [22]: depths, scores = [], []



```

for max_depth in range(1, 20):
    clf = RandomForestClassifier(max_depth=max_depth)
    score = cross_val_score(clf, x1.loc[:50000], y_train.loc[:50000, "ReturnQuantityBin"], cv=5).mean()
    depths.append(max_depth)
    scores.append(score)
print 'max_depth:', max_depth, '\tscore:', score

```

```


max_depth: 1      score: 0.7912241767474417
max_depth: 2      score: 0.7912241767474417
max_depth: 3      score: 0.7912241767474417
max_depth: 4      score: 0.7912241767474417
max_depth: 5      score: 0.7906641747476417
max_depth: 6      score: 0.7911041787476417
max_depth: 7      score: 0.7913441707488417
max_depth: 8      score: 0.7910241787496418
max_depth: 9      score: 0.790704192744642
max_depth: 10     score: 0.7899842167458422
max_depth: 11     score: 0.7900041887442419
max_depth: 12     score: 0.7899841767450418
max_depth: 13     score: 0.7900842087410421
max_depth: 14     score: 0.7891041727390418
max_depth: 15     score: 0.7884041587376416
max_depth: 16     score: 0.7885242147336422
max_depth: 17     score: 0.7860241727158417
max_depth: 18     score: 0.7870443087298431
max_depth: 19     score: 0.7839042567114426

```

```

In [23]: plt.figure(figsize=(16, 10))
         plt.xlabel("depth")
         plt.ylabel("test score")
         plt.plot(depths, scores, '-+')
         plt.show()

```



Another attempt, without cross validation, just we dataset shuffling.

```
In [19]: for depth in range(1, 20):  
          scores_train, scores_test, best_clf, best_score = shuffle(x1, y_train, lambda : RandomForestClassifier(max_depth=depth), step  
          print "depth:", depth, "\tscore:", best_score
```


```
depth: 1          score: 0.6127363618489338
```

depth: 2	score: 0.6189378866971272
depth: 3	score: 0.6321598287938053
depth: 4	score: 0.6403760642500874
depth: 5	score: 0.6412143144576907
depth: 6	score: 0.6514337313548562
depth: 7	score: 0.6432620424000073
depth: 8	score: 0.6457797828784382
depth: 9	score: 0.6400541246853311
depth: 10	score: 0.6460751818421664
depth: 11	score: 0.6393927612274345
depth: 12	score: 0.6396849296339291
depth: 13	score: 0.6330789994311349
depth: 14	score: 0.6318776179566811
depth: 15	score: 0.6240696201724464
depth: 16	score: 0.6212071705903822
depth: 17	score: 0.6120505879455943
depth: 18	score: 0.6141664183891156
depth: 19	score: 0.6093554898816067

Now we try

```
In [20]: shuffle(x1, y_train, lambda : RandomForestClassifier(max_depth=6), steps=10)
```

test 0	train: 0.6701269236262677	test: 0.6477319675303895
test 1	train: 0.6714602896444263	test: 0.6461781632003937
test 2	train: 0.6648065317801488	test: 0.6458100691117841
test 3	train: 0.6647409730713175	test: 0.6400528431178711
test 4	train: 0.6731866436449108	test: 0.63990671741453
test 5	train: 0.6718527999500885	test: 0.6424648218756919
test 6	train: 0.6646784939167376	test: 0.6374748539568927
test 7	train: 0.673734724850982	test: 0.6394368550570327
test 8	train: 0.6641498890165085	test: 0.6428375878644402
test 9	train: 0.6655171647298975	test: 0.6460130882283992



```
Out[20]: ([0.6701269236262677,  
           0.6714602896444263,  
           0.6648065317801488,  
           0.6647409730713175,  
           0.6731866436449108,
```

```

0.6718527999500885,
0.6646784939167376,
0.673734724850982,
0.6641498890165085,
0.6655171647298975],
[0.6477319675303895,
0.6461781632003937,
0.6458100691117841,
0.6400528431178711,
0.63990671741453,
0.6424648218756919,
0.6374748539568927,
0.6394368550570327,
0.6428375878644402,
0.6460130882283992],
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=6, max_features='auto', max_leaf_nodes=None,
                        min_impurity_split=1e-07, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=10, n_jobs=1, oob_score=False, random_state=None,
                        verbose=0, warm_start=False),
0.6477319675303895)

```

```

In [169]: def prune(x1, x2, y, threshold=1, slc=100000, max_depth=9, do_score=True, do_print=True, do_plot=True):
    # building the random forest classifier on a suffled training set
    slc = 100000
    z = x1.copy(deep=True)
    z["ReturnQuantityBin"] = y_train.ReturnQuantityBin
    u = z.sample(frac=1)
    v = u.loc[:, ["ReturnQuantityBin"]]
    u = u.drop(["ReturnQuantityBin"], axis=1)
    forest = RandomForestClassifier(max_depth=9)
    forest.fit(u.iloc[:slc], v.ReturnQuantityBin[:slc])

    # computing test scores
    if do_score:
        predict_test = forest.predict_proba(u.iloc[slc:2 * slc])
        score_test = roc_auc_score(v.ReturnQuantityBin[slc:2 * slc], predict_test[:, 1])
        print "test score:", score_test

```

```

importances = forest.feature_importances_
std = np.std([tree.feature_importances_ for tree in forest.estimators_], axis=0)
indices = np.argsort(importances)[::-1]

# print the feature ranking
if do_print:
    print "Feature ranking:"
    for f in range(u.shape[1]):
        print("%d. %s (%f)" % (f + 1, u.columns[indices[f]], importances[indices[f]]))

# plot the feature ranking
if do_plot:
    plt.figure(figsize=(16, 10))
    plt.title("Feature importances")
    plt.bar(range(u.shape[1]), importances[indices], color="r", yerr=std[indices], align="center")
    plt.xticks(range(u.shape[1]), indices)
    plt.xlim([-1, u.shape[1]])
    plt.show()

# pruning x1 and x2
whitelist = [u.columns[indices[f]] for f in range(min(threshold, len(u.columns)))]
x3 = x1.drop([col for col in x1.columns if col not in whitelist], axis=1)
x4 = x2.drop([col for col in x2.columns if col not in whitelist], axis=1)
return x3, x4

def find_threshold(x1, x2, y, slc=100000, max_depth=9):
    # building the random forest classifier on a suffled training set
    slc = 100000
    z = x1.copy(deep=True)
    z["ReturnQuantityBin"] = y_train.ReturnQuantityBin
    u = z.sample(frac=1)
    v = u.loc[:, ["ReturnQuantityBin"]]
    u = u.drop(["ReturnQuantityBin"], axis=1)
    forest = RandomForestClassifier(max_depth=9)
    forest.fit(u.iloc[:slc], v.ReturnQuantityBin[:slc])

    importances = forest.feature_importances_
    std = np.std([tree.feature_importances_ for tree in forest.estimators_], axis=0)

```

```

indices = np.argsort(importances)[::-1]

best_t, best_score = None, None
for threshold in range(1, 101):
    whitelist = [u.columns[indices[f]] for f in range(min(threshold, len(u.columns)))]
    x3 = x1.drop([col for col in x1.columns if col not in whitelist], axis=1)
    x4 = x2.drop([col for col in x2.columns if col not in whitelist], axis=1)
    sc_train, sc_test, clf, score = shuffle(x3, y_train, slc=100000, steps=2, plot=False)
    if best_t is None or best_score < score:
        best_t, best_score = threshold, score
    print "threshold:", threshold, "\tshape:", x3.shape, "\tscore:", score

return best_t, best_score

```

In [170]: find\_threshold(x1, x2, y\_train)

threshold: 1	shape: (1067290, 1)	score: 0.574356605386431
threshold: 2	shape: (1067290, 2)	score: 0.5856617291218396
threshold: 3	shape: (1067290, 3)	score: 0.6154359104170999
threshold: 4	shape: (1067290, 4)	score: 0.6337566424354024
threshold: 5	shape: (1067290, 5)	score: 0.6351526668057986
threshold: 6	shape: (1067290, 6)	score: 0.6335189564429284
threshold: 7	shape: (1067290, 7)	score: 0.6355780083385949
threshold: 8	shape: (1067290, 8)	score: 0.636939926419611
threshold: 9	shape: (1067290, 9)	score: 0.6358006403113113
threshold: 10	shape: (1067290, 10)	score: 0.6405117479036697
threshold: 11	shape: (1067290, 11)	score: 0.6415424440098114
threshold: 12	shape: (1067290, 12)	score: 0.6412234090026194
threshold: 13	shape: (1067290, 13)	score: 0.6443138341425757
threshold: 14	shape: (1067290, 14)	score: 0.6420528173621186
threshold: 15	shape: (1067290, 15)	score: 0.6456081277227377
threshold: 16	shape: (1067290, 16)	score: 0.6447569832212464
threshold: 17	shape: (1067290, 17)	score: 0.6457928198211941
threshold: 18	shape: (1067290, 18)	score: 0.6463309977585872
threshold: 19	shape: (1067290, 19)	score: 0.6469198407627899
threshold: 20	shape: (1067290, 20)	score: 0.6459300320269035
threshold: 21	shape: (1067290, 21)	score: 0.6451033523300391
threshold: 22	shape: (1067290, 22)	score: 0.646443826512408
threshold: 23	shape: (1067290, 23)	score: 0.6476287452149596
threshold: 24	shape: (1067290, 24)	score: 0.6489343347297963

threshold: 25	shape: (1067290, 25)	score: 0.6487964287409811
threshold: 26	shape: (1067290, 26)	score: 0.6448783857278018
threshold: 27	shape: (1067290, 27)	score: 0.6485368465116832
threshold: 28	shape: (1067290, 28)	score: 0.6486301583341081
threshold: 29	shape: (1067290, 29)	score: 0.6485144251326611
threshold: 30	shape: (1067290, 30)	score: 0.64890651334502
threshold: 31	shape: (1067290, 31)	score: 0.6477219554626193
threshold: 32	shape: (1067290, 32)	score: 0.648766637340102
threshold: 33	shape: (1067290, 33)	score: 0.6532069089350414
threshold: 34	shape: (1067290, 34)	score: 0.6522325246425472
threshold: 35	shape: (1067290, 35)	score: 0.6504259565839989
threshold: 36	shape: (1067290, 36)	score: 0.6488052752460491
threshold: 37	shape: (1067290, 37)	score: 0.6539173510060831
threshold: 38	shape: (1067290, 38)	score: 0.6491065313888993
threshold: 39	shape: (1067290, 39)	score: 0.6502308793567172
threshold: 40	shape: (1067290, 40)	score: 0.6506077866574398
threshold: 41	shape: (1067290, 41)	score: 0.6537339130898341
threshold: 42	shape: (1067290, 42)	score: 0.6510050152717761
threshold: 43	shape: (1067290, 43)	score: 0.6530424101535781
threshold: 44	shape: (1067290, 44)	score: 0.6474548756821336
threshold: 45	shape: (1067290, 45)	score: 0.6507491900527157
threshold: 46	shape: (1067290, 46)	score: 0.653299574890476
threshold: 47	shape: (1067290, 47)	score: 0.6529056725471153
threshold: 48	shape: (1067290, 48)	score: 0.6523421940365526
threshold: 49	shape: (1067290, 49)	score: 0.6524679795342836
threshold: 50	shape: (1067290, 50)	score: 0.6482041290684625
threshold: 51	shape: (1067290, 51)	score: 0.6539802570402597
threshold: 52	shape: (1067290, 52)	score: 0.6517596281633151
threshold: 53	shape: (1067290, 53)	score: 0.6525032299485756
threshold: 54	shape: (1067290, 54)	score: 0.6518686850616213
threshold: 55	shape: (1067290, 55)	score: 0.6533365441127534
threshold: 56	shape: (1067290, 56)	score: 0.6504372939996386
threshold: 57	shape: (1067290, 57)	score: 0.653757515831481
threshold: 58	shape: (1067290, 58)	score: 0.6521883021129744
threshold: 59	shape: (1067290, 59)	score: 0.6531565501152934
threshold: 60	shape: (1067290, 60)	score: 0.6512906521925255
threshold: 61	shape: (1067290, 61)	score: 0.6534267007440926
threshold: 62	shape: (1067290, 62)	score: 0.6526592756398673
threshold: 63	shape: (1067290, 63)	score: 0.6559217414276409



threshold: 64	shape: (1067290, 64)	score: 0.6560805199436732
threshold: 65	shape: (1067290, 65)	score: 0.651927870894848
threshold: 66	shape: (1067290, 66)	score: 0.6541558660710156
threshold: 67	shape: (1067290, 67)	score: 0.6546857848590313
threshold: 68	shape: (1067290, 68)	score: 0.657293670948485
threshold: 69	shape: (1067290, 69)	score: 0.6550863599736362
threshold: 70	shape: (1067290, 70)	score: 0.6517567119092338
threshold: 71	shape: (1067290, 71)	score: 0.6535258569226026
threshold: 72	shape: (1067290, 72)	score: 0.6537079053854153
threshold: 73	shape: (1067290, 73)	score: 0.6519155054612111
threshold: 74	shape: (1067290, 74)	score: 0.6536214599344954
threshold: 75	shape: (1067290, 75)	score: 0.6534983552654517
threshold: 76	shape: (1067290, 76)	score: 0.6548824344874385
threshold: 77	shape: (1067290, 77)	score: 0.6536674290767304
threshold: 78	shape: (1067290, 78)	score: 0.6522533658478034
threshold: 79	shape: (1067290, 79)	score: 0.6533295865622228
threshold: 80	shape: (1067290, 80)	score: 0.6531193399911166
threshold: 81	shape: (1067290, 81)	score: 0.6487721276743434
threshold: 82	shape: (1067290, 82)	score: 0.6526855961912508
threshold: 83	shape: (1067290, 83)	score: 0.6546857501197395
threshold: 84	shape: (1067290, 84)	score: 0.6539547212741288
threshold: 85	shape: (1067290, 85)	score: 0.6545817634214093
threshold: 86	shape: (1067290, 86)	score: 0.6543700013215863
threshold: 87	shape: (1067290, 87)	score: 0.6520411843948619
threshold: 88	shape: (1067290, 88)	score: 0.6560924354716453
threshold: 89	shape: (1067290, 89)	score: 0.6529521079284721
threshold: 90	shape: (1067290, 90)	score: 0.6529490199540005
threshold: 91	shape: (1067290, 91)	score: 0.6554780721719459
threshold: 92	shape: (1067290, 92)	score: 0.655081067871826
threshold: 93	shape: (1067290, 93)	score: 0.6519487659659278
threshold: 94	shape: (1067290, 94)	score: 0.6541754728708041
threshold: 95	shape: (1067290, 95)	score: 0.6514389515947504
threshold: 96	shape: (1067290, 96)	score: 0.6534321968221015
threshold: 97	shape: (1067290, 97)	score: 0.6517552319654643
threshold: 98	shape: (1067290, 98)	score: 0.655481856314928
threshold: 99	shape: (1067290, 99)	score: 0.6546152950705799
threshold: 100	shape: (1067290, 100)	score: 0.6571203334358342

```
Out[170]: (68, 0.657293670948485)
```

```
In [158]: x3, x4 = prune(x1, x2, y_train, threshold=68)
```

test 0	train: 0.6459979624375067	test: 0.6465171859759469
test 1	train: 0.6469237937231148	test: 0.6502324199041669
test 2	train: 0.6456313523173876	test: 0.6484931902863063
test 3	train: 0.6485272949490786	test: 0.646645495718916
test 4	train: 0.6466164986683063	test: 0.6444993330427363
test 5	train: 0.6475535061472868	test: 0.6456752369675604
test 6	train: 0.6459622359306412	test: 0.6501444747268446
test 7	train: 0.6469440086651036	test: 0.6460775489289443
test 8	train: 0.648579527639573	test: 0.6451953263351881
test 9	train: 0.6469833710062032	test: 0.6441433060618417

